

# Introduction

## Deep Learning

*Deep Neural Networks*

*Deep Structural Learning*

*Deep Belief Networks*

# DL is providing breakthrough results in speech recognition and image classification ...

From this Hinton et al 2012 paper:

<http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/38131.pdf>

modeling technique	#params [10 <sup>6</sup> ]	WER		task	hours of training data	DNN-HMM	GMM-HMM with same data	GMM-HMM with more data
		Hub5'00-SWB	RT03S-FSH					
GMM, 40 mix DT 309h SI	29.4	23.6	27.4	Switchboard (test set 1)	309	18.5	27.4	18.6 (2000 hrs)
				Switchboard (test set 2)	309	16.1	23.6	17.1 (2000 hrs)
NN 1 hidden-layer×4634 units	43.6	26.0	29.4	English Broadcast News	50	17.5	18.8	
+ 2×5 neighboring frames	45.1	22.4	25.7	Bing Voice Search	24	30.4	36.2	
DBN-DNN 7 hidden layers×2048 units	45.1	17.1	19.6	(Sentence error rates)				
+ updated state alignment	45.1	16.4	18.6	Google Voice Input	5,870	12.3		16.0 (>>5,870hrs)
+ sparsification	15.2 nz	16.1	18.5	Youtube	1,400	47.6	52.3	
GMM 72 mix DT 2000h SA	102.4	17.1	18.6					

go here: <http://yann.lecun.com/exdb/mnist/>

From here:

<http://people.idsia.ch/~juergen/cvpr2012.pdf>

Dataset	Best result of others [%]	MCDNN [%]	Relative improv. [%]
MNIST	0.39	0.23	41
NIST SD 19	see Table 4	see Table 4	30-80
HWDB1.0 on.	7.61	5.61	26
HWDB1.0 off.	10.01	6.5	35
CIFAR10	18.50	11.21	39
traffic signs	1.69	0.54	72
NORB	5.00	2.70	46

So, 1. **what exactly is deep learning ?**

And, 2. **why is it generally better** than other methods on image, speech and certain other types of data?

So, 1. **what exactly is deep learning ?**

And, 2. **why is it generally better** than other methods on image, speech and certain other types of data?

### **The short answers**

- 1. ‘Deep Learning’ means using a neural network with several layers of nodes between input and output**
- 2. the series of layers between input & output do feature identification and processing in a series of stages, just as our brains seem to.**


hmmm... OK, but:

**3. multilayer neural networks have been around for 25 years. What's actually new?**


hmmm... OK, but:

**3. multilayer neural networks have been around for 25 years. What's actually new?**

**we have always had good algorithms for learning the weights in networks with 1 hidden layer**



**but these algorithms are not good at learning the weights for networks with more hidden layers**




**what's new is: algorithms for training many-layer networks**

# longer answers

1. reminder/quick-explanation of how neural network weights are learned;
2. the idea of **unsupervised feature learning** (why ‘intermediate features’ are important for difficult classification tasks, and how NNs seem to naturally learn them)
3. The ‘breakthrough’ – the simple trick for training Deep neural networks

$$f(x) = \frac{1}{1 + e^{-x}}$$



-0.06

W1

-2.5

W2


W3

1.4

$f(x)$



$$f(x) = \frac{1}{1 + e^{-x}}$$



-0.06

2.7

-2.5

-8.6

0.002

1.4

$f(x)$

$$x = -0.06 \times 2.7 + 2.5 \times 8.6 + 1.4 \times 0.002 = 21.34$$

*A dataset*

<i>Fields</i>	<i>class</i>
---------------	--------------


1.4 2.7 1.9	0
-------------	---

3.8 3.4 3.2	0
-------------	---

6.4 2.8 1.7	1
-------------	---

4.1 0.1 0.2	0
-------------	---

etc ...



## *Training the neural network*

***Fields***                      ***class***


1.4 2.7 1.9                0

3.8 3.4 3.2                0

6.4 2.8 1.7                1

4.1 0.1 0.2                0

etc ...



*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Initialise with random weights**



*Training data*

*Fields* *class*

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Present a training pattern



*Training data*

*Fields* *class*

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Feed it through to get output



*Training data*

*Fields* *class*

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Compare with target output



*Training data*

*Fields* *class*

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Adjust weights based on error





*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Present a training pattern**



*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Feed it through to get output**



*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Compare with target output**



*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Adjust weights based on error**



*Training data*

***Fields***                      ***class***

1.4 2.7 1.9              0


3.8 3.4 3.2              0

6.4 2.8 1.7              1

4.1 0.1 0.2              0

etc ...

And so on ....




Repeat this thousands, maybe millions of times – each time taking a random training instance, and making slight weight adjustments

*Algorithms for weight adjustment are designed to make changes that will reduce the error*


# The decision boundary perspective...

Initial random weights




# The decision boundary perspective...

Present a training instance / adjust the weights



# The decision boundary perspective...


Present a training instance / adjust the weights






# The decision boundary perspective...

Present a training instance / adjust the weights



# The decision boundary perspective...

Present a training instance / adjust the weights




# The decision boundary perspective...

Eventually ....



# The point is that

- weight-learning algorithms for NNs are dumb
- they work by making thousands and thousands of tiny adjustments, each making the network do better at the most recent pattern, but perhaps a little worse on many others
- but, by dumb luck, eventually this tends to be good enough to learn effective classifiers for many real applications




# Some other points

**Detail** of a standard NN weight learning algorithm – **previously introduced**

If  $f(x)$  is non-linear, a network with 1 hidden layer can, in theory, learn perfectly any classification problem. A set of weights exists that can produce the targets from the inputs. The problem is finding them.


# Some other ‘by the way’ points

If  $f(x)$  is linear, the NN can **only** draw straight decision boundaries (even if there are many layers of units)




# Some other ‘by the way’ points

NNs use nonlinear  $f(x)$  so they  
can draw complex boundaries,  
but keep the data unchanged




# Some other ‘by the way’ points


NNs use nonlinear  $f(x)$  so they can draw complex boundaries, but keep the data unchanged



SVMs only draw straight lines, but they transform the data first in a way that makes that OK










# Feature detectors



Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*






*what is this  
unit doing?*


Figure 1.2: Examples of handwritten digits from U.S. postal envelopes.




# Hidden layer units become *self-organised feature detectors*




# What does this unit detect?




# What does this unit detect?



# What does this unit detect?



# What does this unit detect?






Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

What features might you expect a good NN to learn, when trained with data like this?






Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*




Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*


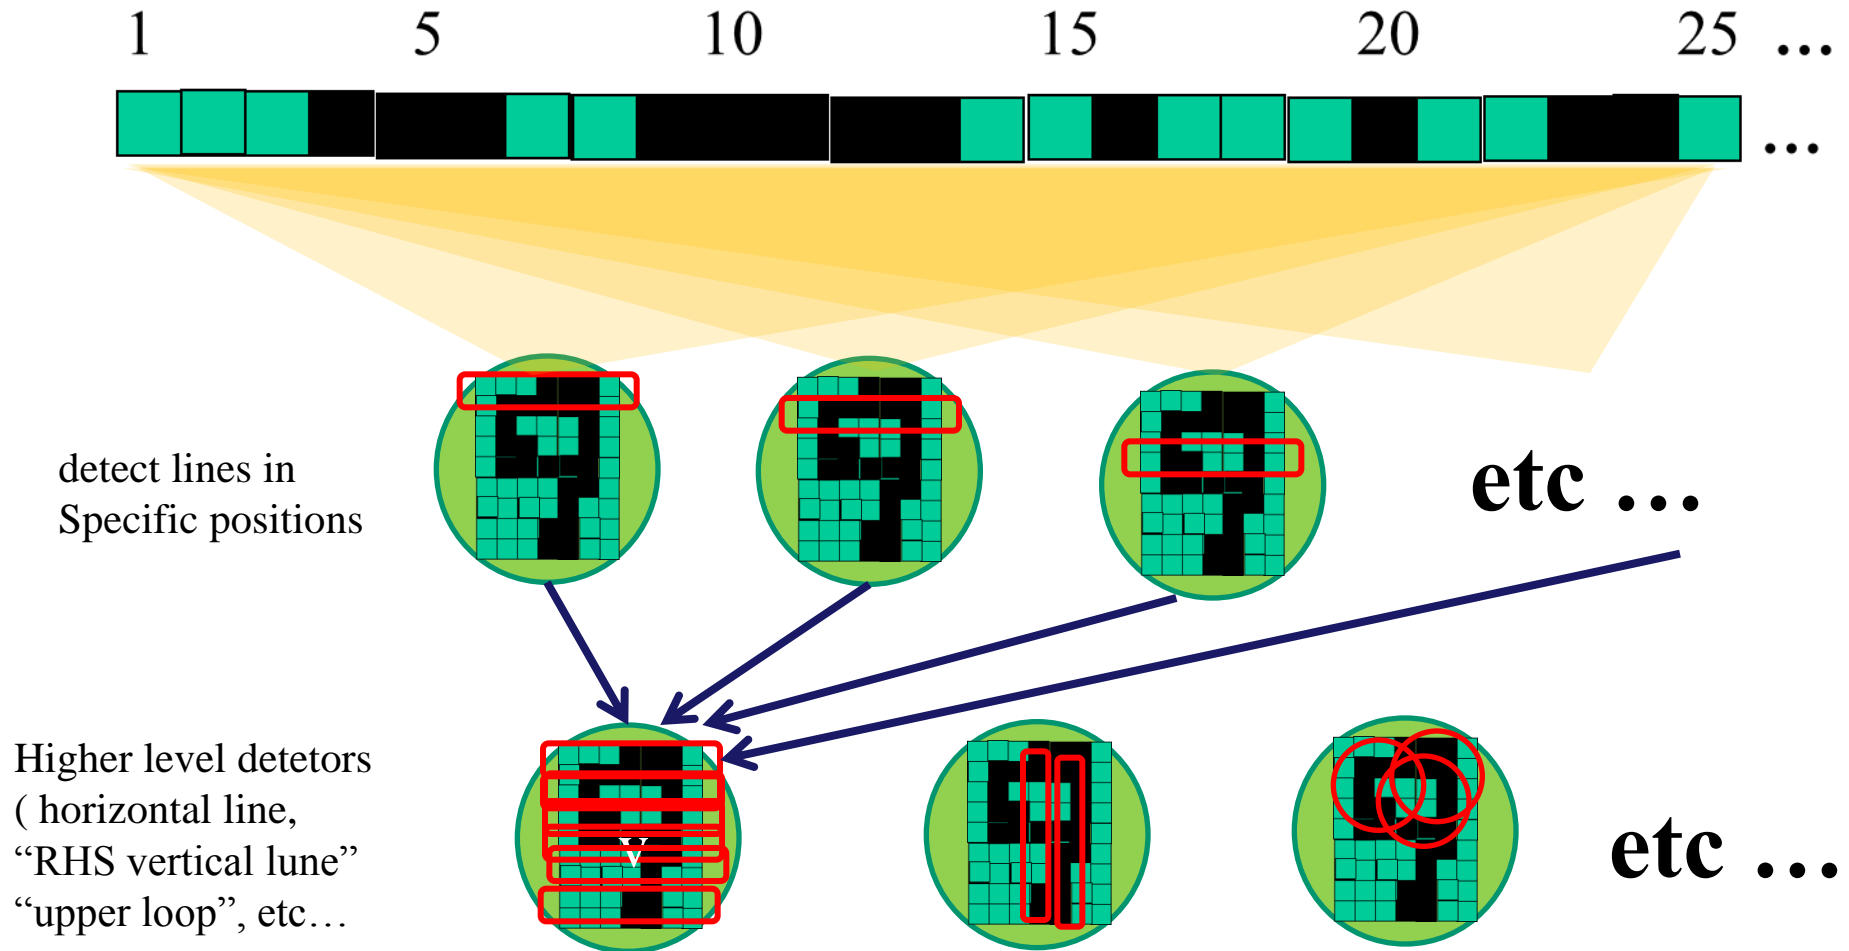


Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

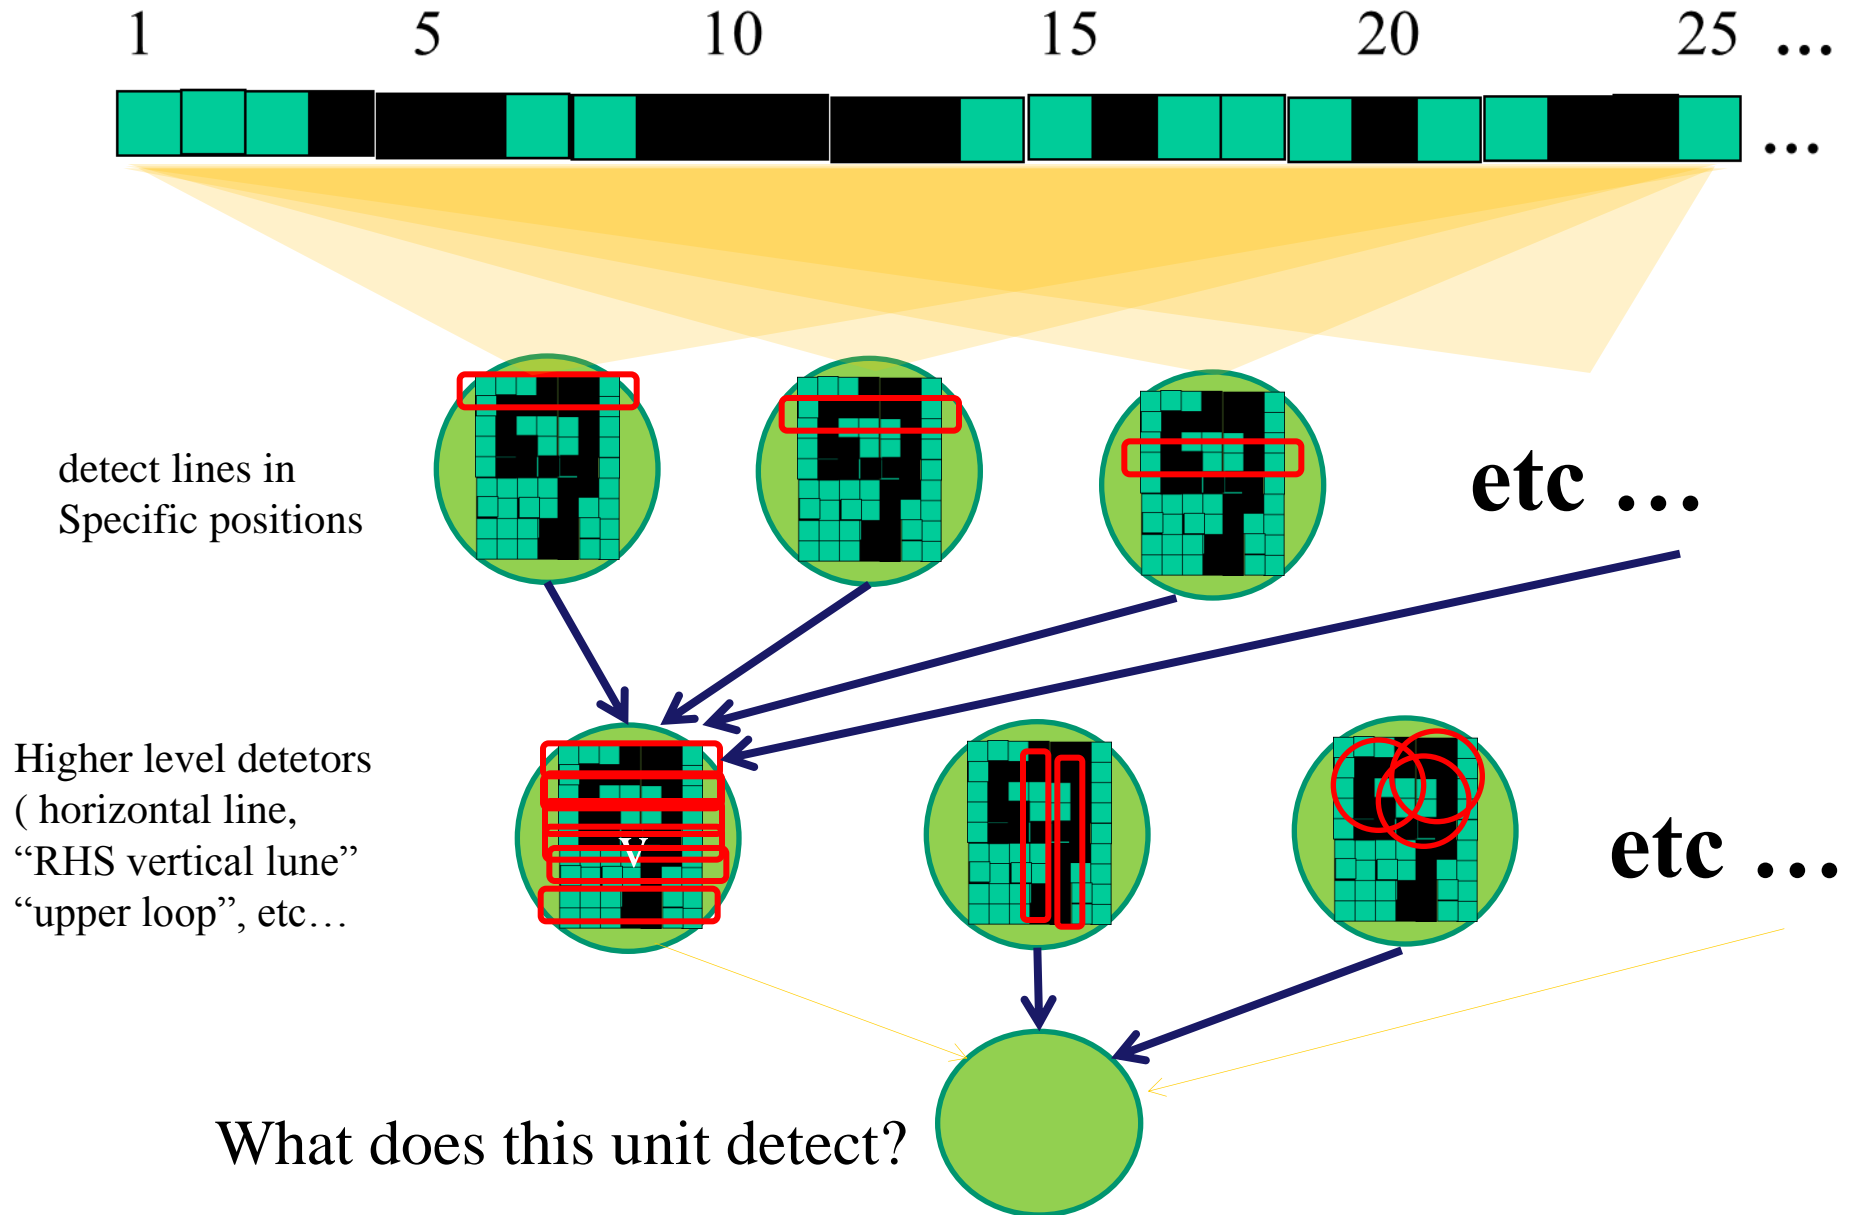


But what about position invariance ???  
our example unit detectors were tied to  
specific parts of the image


successive layers can learn higher-level features ...



successive layers can learn higher-level features ...





*So: multiple layers make sense*



# *So: multiple layers make sense*


## Your brain works that way







# *So: multiple layers make sense*

**Many-layer neural network architectures should be capable of learning the true underlying features and ‘feature logic’, and therefore generalise very well ...**




But, until very recently, our weight-learning algorithms simply did not work on multi-layer architectures




Along came deep learning ...

# The new way to train multi-layer NNs...




# The new way to train multi-layer NNs...



Train **this** layer first


# The new way to train multi-layer NNs...



Train **this** layer first

then **this** layer

# The new way to train multi-layer NNs...




Train **this** layer first

then **this** layer

then **this** layer

# The new way to train multi-layer NNs...



Train **this** layer first


then **this** layer

then **this** layer

then **this** layer



# The new way to train multi-layer NNs...



Train **this** layer first


then **this** layer

then **this** layer

then **this** layer

finally **this** layer


# The new way to train multi-layer NNs...




*EACH of the (non-output) layers is trained to be an **auto-encoder***

*Basically, it is forced to learn good features that describe what comes from the previous layer*

**an auto-encoder is trained, with an absolutely standard weight-adjustment algorithm to reproduce the input**




**an auto-encoder is trained, with an absolutely standard weight-adjustment algorithm to reproduce the input**




**By making this happen with (many) fewer units than the inputs, this forces the ‘hidden layer’ units to become good feature detectors**

intermediate layers are each trained to be auto encoders (or similar)



Final layer trained to predict class based on outputs from previous layers



# And that's that

- That's the basic idea
- There are many many types of deep learning,
- different kinds of autoencoder, variations on architectures and training algorithms, etc...
- Very fast growing area ...

